

CSL7670 : Fundamentals of Machine Learning

Lab Report



॥ त्वं ज्ञानमयो विज्ञानमयोऽसि ॥

Name:	Ganesh Kumar Nagal
Roll Number:	M23MEA004
Program:	M.Tech Advance Manufacturing And Design

Chapter 1

Lab-4

1.1 Objective

Objective of this assignment is to gain familiarity with neural networks (multi-layer perception).

1.2 Problem-1

(Simple MLP) Please go through this blog on developing your first neural network: [Blog] and understand the code.

1. (A) Draw the model architecture by showing each perceptron, input/output/hidden layers.
2. (B) Change the code to use only 60%, 70%, and 80% data as training, and report test-set performance for all these three training data set sizes.
3. (C) Compare BCE loss with MSE loss.
4. (D) Change the number of hidden layers from 2 to 4 and compare the performance.

Solution 1:

```
1  # -*- coding: utf-8 -*-
2  """Untitled5.ipynb
3
4  Automatically generated by Colaboratory.
5
6  Original file is located at
7      https://colab.research.google.com/drive/1
      ↪ RPbrLSJkg9r1bPcAGKDAHRcbmkg8kU0
8  """
9
10 from google.colab import files
11 files = files.upload()
12
13 #sir code
14 import numpy as np
15 import torch
16 import torch.nn as nn
17 import torch.optim as optim
18
19 # load the dataset, split into input (X) and output (y) variables
20 dataset = np.loadtxt('pima-indians-diabetes.csv', delimiter=',')
```

```

21 X = dataset[:,0:8]
22 y = dataset[:,8]
23
24 X = torch.tensor(X, dtype=torch.float32)
25 y = torch.tensor(y, dtype=torch.float32).reshape(-1, 1)
26
27 # define the model
28 class PimaClassifier(nn.Module):
29     def __init__(self):
30         super().__init__()
31         self.hidden1 = nn.Linear(8, 12)
32         self.act1 = nn.ReLU()
33         self.hidden2 = nn.Linear(12, 8)
34         self.act2 = nn.ReLU()
35         self.output = nn.Linear(8, 1)
36         self.act_output = nn.Sigmoid()
37
38     def forward(self, x):
39         x = self.act1(self.hidden1(x))
40         x = self.act2(self.hidden2(x))
41         x = self.act_output(self.output(x))
42         return x
43
44 model = PimaClassifier()
45 print(model)
46
47 # train the model
48 loss_fn = nn.BCELoss() # binary cross entropy
49 optimizer = optim.Adam(model.parameters(), lr=0.001)
50
51 n_epochs = 100
52 batch_size = 10
53
54 for epoch in range(n_epochs):
55     for i in range(0, len(X), batch_size):
56         Xbatch = X[i:i+batch_size]
57         y_pred = model(Xbatch)
58         ybatch = y[i:i+batch_size]
59         loss = loss_fn(y_pred, ybatch)
60         optimizer.zero_grad()
61         loss.backward()
62         optimizer.step()
63
64 # compute accuracy
65 y_pred = model(X)
66 accuracy = (y_pred.round() == y).float().mean()
67 print(f"Accuracy_{accuracy}")
68
69 # make class predictions with the model
70 predictions = (model(X) > 0.5).int()
71 for i in range(5):
72     print('%s=>%d_(expected_%d)' % (X[i].tolist(), predictions[i], y[i])
73         ↪ )

```

```

74 #q1 b
75 import numpy as np
76 import torch
77 import torch.nn as nn
78 import torch.optim as optim
79
80 # Load the dataset, split into input (X) and output (y) variables
81 dataset = np.loadtxt('pima-indians-diabetes.csv', delimiter=',')
82 X = dataset[:, 0:8]
83 y = dataset[:, 8]
84
85 X = torch.tensor(X, dtype=torch.float32)
86 y = torch.tensor(y, dtype=torch.float32).reshape(-1, 1)
87
88 # Iterate over different percentages (60%, 70%, and 80%)
89 for num in [0.6, 0.7, 0.8]:
90     full_data = len(X)
91     train_size = int(num* full_data)
92
93     X_train = X[:train_size]
94     y_train = y[:train_size]
95     X_test = X[train_size:]
96     y_test = y[train_size:]
97
98 # Define the model
99 class PimaClassifier(nn.Module):
100     def __init__(self):
101         super().__init__()
102         self.hidden1 = nn.Linear(8, 12)
103         self.act1 = nn.ReLU()
104         self.hidden2 = nn.Linear(12, 8)
105         self.act2 = nn.ReLU()
106         self.output = nn.Linear(8, 1)
107         self.act_output = nn.Sigmoid()
108
109     def forward(self, x):
110         x = self.act1(self.hidden1(x))
111         x = self.act2(self.hidden2(x))
112         x = self.act_output(self.output(x))
113         return x
114
115 model = PimaClassifier()
116 #print(model)
117
118 # Train the model
119 loss_fn = nn.BCELoss() # Binary Cross-Entropy
120 optimizer = optim.Adam(model.parameters(), lr=0.001)
121
122 n_epochs = 100
123 batch_size = 10
124
125 for epoch in range(n_epochs):
126     for i in range(0, train_size, batch_size):
127         Xbatch = X_train[i:i+batch_size]

```

```

128         y_pred = model(Xbatch)
129         ybatch = y_train[i:i+batch_size]
130         loss = loss_fn(y_pred, ybatch)
131         optimizer.zero_grad()
132         loss.backward()
133         optimizer.step()
134
135         # Compute accuracy on the test set
136         y_pred_test = model(X_test)
137         accuracy = (y_pred_test.round() == y_test).float().mean()
138         print(f"Accuracy for {int(num*100)}% of the dataset as Test Data is:
           ↪ {accuracy}")
139
140 #q1 c
141 from google.colab import files
142 files = files.upload()
143 import numpy as np
144 import torch
145 import torch.nn as nn
146 import torch.optim as optim
147
148 dataset = np.loadtxt('pima-indians-diabetes.csv', delimiter=',')
149 X = dataset[:, 0:8]
150 y = dataset[:, 8]
151
152 X = torch.tensor(X, dtype=torch.float32)
153 y = torch.tensor(y, dtype=torch.float32).reshape(-1, 1)
154
155 class PimaClassifier(nn.Module):
156     def __init__(self):
157         super().__init__()
158         self.hidden1 = nn.Linear(8, 12)
159         self.act1 = nn.ReLU()
160         self.hidden2 = nn.Linear(12, 8)
161         self.act2 = nn.ReLU()
162         self.output = nn.Linear(8, 1)
163         self.act_output = nn.Sigmoid()
164
165     def forward(self, x):
166         x = self.act1(self.hidden1(x))
167         x = self.act2(self.hidden2(x))
168         x = self.act_output(self.output(x))
169         return x
170
171 model = PimaClassifier()
172 print(model)
173
174 # Step 1: Train the model with BCE loss with learning rate 0.001
175 bce_loss_function = nn.BCELoss()
176 bce_optimizer = optim.Adam(model.parameters(), lr=0.001)
177
178 #Step 2: Train the model with MSE loss with learning rate 0.001
179 mse_loss_function = nn.MSELoss()
180 mse_optimizer = optim.Adam(model.parameters(), lr=0.001)

```

```

181
182 n_epochs = 100
183 batch_size = 10
184
185 for epoch in range(n_epochs):
186     for i in range(0, len(X), batch_size):
187         Xbatch = X[i:i+batch_size]
188         ybatch = y[i:i+batch_size]
189
190         # Train with BCE loss
191         bce_optimizer.zero_grad()
192         y_pred_bce = model(Xbatch)
193         bce_loss = bce_loss_function(y_pred_bce, ybatch)
194         bce_loss.backward()
195         bce_optimizer.step()
196
197         # Train with MSE loss
198         mse_optimizer.zero_grad()
199         y_pred_mse = model(Xbatch)
200         mse_loss = mse_loss_function(y_pred_mse, ybatch)
201         mse_loss.backward()
202         mse_optimizer.step()
203
204 # Compute accuracy with BCE loss
205 y_pred_bce = model(X)
206 accuracy_bce = ((y_pred_bce.round() == y).float().mean()).item()
207 print(f"Accuracy with BCE Loss: {accuracy_bce}")
208
209 # Compute accuracy with MSE loss
210 accuracy_mse = ((y_pred_mse.round() == ybatch).float().mean()).item()
211 print(f"Accuracy with MSE Loss: {accuracy_mse}")
212
213 # Print final BCE and MSE losses
214 print(f"Final BCE Loss: {bce_loss.item()}")
215 print(f"Final MSE Loss: {mse_loss.item()}")

```

1.3 Problem-2

Develop a neural network that works for MNIST handwritten digit classification.

1. (A) Draw the model architecture.
2. (B) Compare its performance with KNN classification.

Solution 2:

```

1 # -*- coding: utf-8 -*-
2 """Assignment4.ipynb
3
4 Automatically generated by Colaboratory.
5
6 Original file is located at
7     https://colab.research.google.com/drive/1
    ↗ aCwkbEDKi_fA0JzMVyK9X27s0TtsRH-3

```

```

8  """
9
10 import tensorflow as tf
11 from tensorflow.keras import layers, models
12 from tensorflow.keras.datasets import mnist
13
14 # Load and preprocess the MNIST dataset
15 (train_images, train_labels), (test_images, test_labels) = mnist.load_data
    ↪ ()
16
17 # Normalize pixel values to be in the range [0, 1]
18 train_images, test_images = train_images / 255.0, test_images / 255.0
19
20 # Define the neural network model
21 model = models.Sequential()
22
23 # Input layer: Flatten the 28x28 images into a vector
24 model.add(layers.Flatten(input_shape=(28, 28)))
25
26 # Hidden layers
27 model.add(layers.Dense(128, activation='relu')) # 128 neurons, ReLU
    ↪ activation
28 model.add(layers.Dropout(0.2)) # Dropout for regularization
29 model.add(layers.Dense(64, activation='relu')) # 64 neurons, ReLU
    ↪ activation
30
31 # Output layer: 10 neurons for 10 classes (digits 0-9) with softmax
    ↪ activation
32 model.add(layers.Dense(10, activation='softmax'))
33
34 # Compile the model
35 model.compile(optimizer='adam',
36               loss='sparse_categorical_crossentropy',
37               metrics=['accuracy'])
38
39 # Train the model
40 model.fit(train_images, train_labels, epochs=5, batch_size=64,
    ↪ validation_split=0.2)
41
42 # Evaluate the model on the test data
43 test_loss, test_acc = model.evaluate(test_images, test_labels)
44 print(f"Test accuracy: {test_acc*100:.2f}%")
45
46 import numpy as np
47 import matplotlib.pyplot as plt
48
49 # Select three random test images
50 num_samples = test_images.shape[0]
51 random_indices = np.random.choice(num_samples, 3, replace=False)
52 selected_images = test_images[random_indices]
53 selected_labels = test_labels[random_indices]
54
55 # Make predictions on the selected test images
56 predictions = model.predict(selected_images)

```



```

57
58 # Display the images and model predictions
59 for i in range(3):
60     plt.figure(figsize=(3, 3))
61     plt.imshow(selected_images[i], cmap='gray')
62     plt.title(f"True_Label:_{selected_labels[i]}\nPredicted_Label:_{np.
        ↳ argmax(predictions[i])}")
63     plt.axis('off')
64     plt.show()
65
66
67
68 """Diagram of the model architecture for the neural network
69
70 Input (28x28)
71 |
72 Flatten
73 |
74 |--- Dense (128, ReLU)
75 |   |
76 |   |--- Dropout (0.2)
77 |   |
78 |--- Dense (64, ReLU)
79 |   |
80 |--- Dense (10, Softmax)
81
82
83 In this diagram:
84
85 Input (28x28)" represents the input layer with 28x28 pixels for each image
86 ↳ .
87 Flatten" is used to flatten the 2D image into a 1D vector.
88 Dense (128, ReLU)" represents the first hidden layer with 128 neurons and
89 ↳ ReLU activation.
90 Dropout (0.2)" is a dropout layer with a dropout rate of 0.2 for
91 ↳ regularization.
92 Dense (64, ReLU)" is the second hidden layer with 64 neurons and ReLU
93 ↳ activation.
94 Dense (10, Softmax)" is the output layer with 10 neurons for classifying
95 ↳ digits 0-9 using softmax activation.
96
97 """
98
99 ""B) Comparing Neural Network Performance with K-Nearest Neighbors (KNN):
100
101 To compare the performance of the neural network with K-Nearest Neighbors
102 ↳ (KNN) classification on the MNIST handwritten digit classification
103 ↳ task, we need to consider various parameters and settings for both
104 ↳ approaches.
105
106 Neural Network (NN) Parameters and Settings:
107
108 Architecture:The neural network architecture consists of an input layer

```

→ (784 neurons), two hidden layers (128 neurons with ReLU activation
 → and 64 neurons with ReLU activation), and an output layer (10
 → neurons with softmax activation).

Training: The model was trained with the following settings:

Loss Function: Sparse Categorical Cross-Entropy

Optimizer: Adam

Number of Epochs: 5

Batch Size: 64

Validation Split: 20% of the training data

Regularization: Dropout with a rate of 0.2 was applied to the second
 → hidden layer for regularization.

K-Nearest Neighbors (KNN) Parameters and Settings:

Number of Neighbors (k): This is a critical hyperparameter for KNN. We
 → would need to experiment with different values of k to find the
 → optimal one.

Distance Metric: The choice of distance metric, such as Euclidean distance
 → or Manhattan distance, can affect KNN's performance.

Data Preprocessing: It's essential to preprocess the data similarly to the
 → neural network approach. We normalized the pixel values to be in
 → the range [0, 1], which is a common preprocessing step for both
 → methods.

Inference Time: KNN doesn't require training, so it's faster at inference.
 → However, the inference time can vary depending on the size of the
 → dataset and the chosen k value.

Memory Consumption: KNN stores the entire training dataset, which can be
 → memory-intensive for large datasets.

Accuracy: The accuracy of KNN can vary significantly based on the choice
 → of hyperparameters and data preprocessing.

Scalability: KNN's scalability can be an issue for large datasets, as it
 → requires computing distances to all training examples.

Comparing the performance of NN and KNN would involve training the KNN
 → model with various values of k and comparing their accuracy on the
 → same MNIST test dataset. Typically, a well-tuned neural network
 → would achieve higher accuracy on MNIST than KNN, but KNN can serve
 → as a baseline or be suitable for simpler classification tasks.

"""

Downloading data from [https://storage.googleapis.com/tensorflow/tf-keras-](https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz)
 → [datasets/mnist.npz](https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz)

11490434/11490434 [=====] - 11s 1us/step


Epoch 1/5

750/750 [=====] - 2s 2ms/step - loss: 0.3715 -
 → accuracy: 0.8899 - val_loss: 0.1776 - val_accuracy: 0.9490


Epoch 2/5

```
6 750/750 [=====] - 1s 1ms/step - loss: 0.1637 -  
  ↳ accuracy: 0.9512 - val_loss: 0.1243 - val_accuracy: 0.9619  
7 Epoch 3/5  
8 750/750 [=====] - 1s 1ms/step - loss: 0.1202 -  
  ↳ accuracy: 0.9630 - val_loss: 0.1042 - val_accuracy: 0.9689  
9 Epoch 4/5  
10 750/750 [=====] - 1s 2ms/step - loss: 0.0988 -  
  ↳ accuracy: 0.9694 - val_loss: 0.0905 - val_accuracy: 0.9723  
11 Epoch 5/5  
12 750/750 [=====] - 1s 2ms/step - loss: 0.0839 -  
  ↳ accuracy: 0.9735 - val_loss: 0.0888 - val_accuracy: 0.9726  
13 313/313 [=====] - 0s 990us/step - loss: 0.0857 -  
  ↳ accuracy: 0.9722  
14 Test accuracy: 97.22%
```


True Label: 1
Predicted Label: 1



True Label: 4
Predicted Label: 4




True Label: 2
Predicted Label: 2

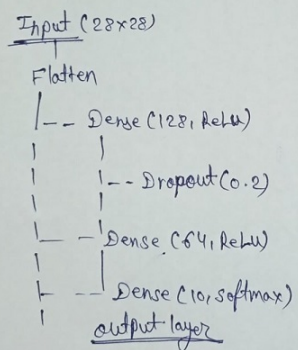


Name - Ganesh Kumar Nagal
Roll no - M23 MEA004

Ans (2) (a) Architecture of model.



Line diagram of Architecture of model



॥ त्वं ज्ञानमयो विद्वानमयोऽसि ॥