

CSL7670 : Fundamentals of Machine Learning

Lab Report



॥ त्वं ज्ञानमयो विज्ञानमयोऽसि ॥

| | |
|--------------|--|
| Name: | Ganesh Kumar Nagal |
| Roll Number: | M23MEA004 |
| Program: | M.Tech Advance Manufacturing And Design |

Chapter 1

Lab-5 and 6 (CNN)

1.1 Objective

Objective of this assignment is to gain familiarity with convolutional neural networks.

1.2 Problem-1

1. (Simple CNN) Go through the following tutorial to understand how to train a CNN classifier: https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html Now,

1. (A) Understand the code completely and run it. perceptron,input/output/hidden layers.
2. (B) (b) Explain the following Pytorch Functions: (i) conv2D (ii) MaxPool2d (iii) Linear (iv) Relu (v) linear.
3. (C) (c) Plot the loss function.
4. (D) (d) Edit the code to modify the CNN architecture in the following four steps (call it myCNN): (i) Instead of 6 activation maps in conv1, use 5 activation maps, (ii) instead of maxpool use average pool, (iii) Instead of 16 activation maps in conv1, use 10 activation maps, and (iv) Remove fc2 and change fc1 so that it projects to 100 dimensions instead of 120 currently. Rerun the experiment and compare CNN (original code) and myCNN (this code).

Solution 1:

```
1 # -*- coding: utf-8 -*-
2 """Untitled3.ipynb
3
4 Automatically generated by Colaboratory.
5
6 Original file is located at
7     https://colab.research.google.com/drive/1f-
8         ↪ CdBABYVKgJKcWzd2XfX2u1p52pcEVb
9
10 1. Load and normalize CIFAR10
11 """
12 import torch
13 import torch.nn as nn
14 import torch.optim as optim
15 import torchvision
16 import torchvision.transforms as transforms
```

```

17 import matplotlib.pyplot as plt
18
19 transform = transforms.Compose(
20     [transforms.ToTensor(),
21      transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
22
23 batch_size = 4
24
25 trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
26                                         download=True, transform=transform
27                                         ↪ )
28 trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
29                                           shuffle=True, num_workers=2)
30
31 testset = torchvision.datasets.CIFAR10(root='./data', train=False,
32                                         download=True, transform=transform)
33 testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
34                                           shuffle=False, num_workers=2)
35
36 classes = ('plane', 'car', 'bird', 'cat',
37            'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
38
39 import matplotlib.pyplot as plt
40 import numpy as np
41
42 # functions to show an image
43
44 def imshow(img):
45     img = img / 2 + 0.5     # unnormalize
46     npimg = img.numpy()
47     plt.imshow(np.transpose(npimg, (1, 2, 0)))
48     plt.show()
49
50
51 # get some random training images
52 dataiter = iter(trainloader)
53 images, labels = next(dataiter)
54
55 # show images
56 imshow(torchvision.utils.make_grid(images))
57 # print labels
58 print(' '.join(f'{classes[labels[j]]:5s}' for j in range(batch_size)))
59
60 """2. Define a Convolutional Neural Network"""
61
62 import torch.nn as nn
63 import torch.nn.functional as F
64
65
66 class Net(nn.Module):
67     def __init__(self):
68         super().__init__()
69         self.conv1 = nn.Conv2d(3, 6, 5)

```

```

70         self.pool = nn.MaxPool2d(2, 2)
71         self.conv2 = nn.Conv2d(6, 16, 5)
72         self.fc1 = nn.Linear(16 * 5 * 5, 120)
73         self.fc2 = nn.Linear(120, 84)
74         self.fc3 = nn.Linear(84, 10)
75
76     def forward(self, x):
77         x = self.pool(F.relu(self.conv1(x)))
78         x = self.pool(F.relu(self.conv2(x)))
79         x = torch.flatten(x, 1) # flatten all dimensions except batch
80         x = F.relu(self.fc1(x))
81         x = F.relu(self.fc2(x))
82         x = self.fc3(x)
83         return x
84
85
86 net = Net()
87
88 """3. Define a Loss function and optimizer"""
89
90 import torch.optim as optim
91
92 criterion = nn.CrossEntropyLoss()
93 optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
94
95 """4. Train the network"""
96
97 # Lists to store training loss values
98 train_losses = []
99
100
101
102 for epoch in range(2): # loop over the dataset multiple times
103
104     running_loss = 0.0
105     for i, data in enumerate(trainloader, 0):
106         # get the inputs; data is a list of [inputs, labels]
107         inputs, labels = data
108
109         # zero the parameter gradients
110         optimizer.zero_grad()
111
112         # forward + backward + optimize
113         outputs = net(inputs)
114         loss = criterion(outputs, labels)
115         loss.backward()
116         optimizer.step()
117
118         # print statistics
119         running_loss += loss.item()
120         if i % 2000 == 1999: # Print every 2000 mini-batches
121             print(f'[{epoch+1}, {i+1:5d}] loss: {running_loss / 2000:.3f}')
122             train_losses.append(running_loss / 2000)

```

```

123         running_loss = 0.0
124
125
126 print('Finished Training')
127
128 """Save trained model"""
129
130 PATH = './cifar_net.pth'
131 torch.save(net.state_dict(), PATH)
132
133 """5. Test the network on the test data"""
134
135 dataiter = iter(testloader)
136 images, labels = next(dataiter)
137
138 # print images
139 imshow(torchvision.utils.make_grid(images))
140 print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] for j in range
    ↪ (4)))
141
142 net = Net()
143 net.load_state_dict(torch.load(PATH))
144
145 # let us see what the neural network thinks these examples above are:
146 outputs = net(images)
147
148 _, predicted = torch.max(outputs, 1)
149
150 print('Predicted: ', ' '.join('%5s' % classes[predicted[j]]
    ↪ for j in range(4)))
151
152
153 # Let us look at how the network performs on the whole dataset.
154 correct = 0
155 total = 0
156 # since we're not training, we don't need to calculate the gradients for
    ↪ our outputs
157 with torch.no_grad():
158     for data in testloader:
159         images, labels = data
160         # calculate outputs by running images through the network
161         outputs = net(images)
162         # the class with the highest energy is what we choose as
            ↪ prediction
163         _, predicted = torch.max(outputs.data, 1)
164         total += labels.size(0)
165         correct += (predicted == labels).sum().item()
166
167 print(f'Accuracy of the network on the 10000 test images: {100 * correct /
    ↪ total} %')
168
169 # the classes that performed well, and the classes that did not perform
    ↪ well:
170
171 # prepare to count predictions for each class

```

```

172 correct_pred = {classname: 0 for classname in classes}
173 total_pred = {classname: 0 for classname in classes}
174
175 # again no gradients needed
176 with torch.no_grad():
177     for data in testloader:
178         images, labels = data
179         outputs = net(images)
180         _, predictions = torch.max(outputs, 1)
181         # collect the correct predictions for each class
182         for label, prediction in zip(labels, predictions):
183             if label == prediction:
184                 correct_pred[classes[label]] += 1
185                 total_pred[classes[label]] += 1
186
187
188 # print accuracy for each class
189 for classname, correct_count in correct_pred.items():
190     accuracy = 100 * float(correct_count) / total_pred[classname]
191     print(f'Accuracy for class: {classname:5s} is {accuracy:.1f}%')
192
193
194
195 """Explanation of following Pytorch Functions:
196
197
198 (i) 'conv2d':
199     - 'conv2d' stands for "convolutional 2D." It is a function in PyTorch
200       ↪ used for 2D convolution operations, which are fundamental in deep
201       ↪ learning for tasks like image processing and computer vision.
202       ↪ Convolution involves applying a filter (also known as a kernel)
203       ↪ to an input image to produce a feature map. The filter slides
204       ↪ over the input, and at each position, it computes a weighted sum
205       ↪ of the input values within its receptive field. This operation is
206       ↪ used to extract features from the input data.
207
208 (ii) 'MaxPool2d':
209     - 'MaxPool2d' is short for "Max Pooling 2D." It is a pooling operation
210       ↪ in PyTorch used primarily in convolutional neural networks (CNNs)
211       ↪ for downsampling and reducing the spatial dimensions of feature
212       ↪ maps. Max pooling works by dividing the input into non-
213       ↪ overlapping regions (typically 2x2 or 3x3), and for each region,
214       ↪ it takes the maximum value. This reduces the size of the feature
215       ↪ maps while retaining the most important information, helping to
216       ↪ reduce computational complexity and prevent overfitting.
217
218 (iii) 'Linear':
219     - 'Linear' is a PyTorch module that represents a fully connected layer
220       ↪ or a linear transformation. In a neural network, this layer
221       ↪ performs a linear mapping of the input data to a set of output
222       ↪ neurons, where each output neuron is connected to every input
223       ↪ neuron. This layer is also known as a dense layer or a fully
224       ↪ connected layer. The linear transformation is typically followed
225       ↪ by an activation function to introduce non-linearity into the

```

→ network.

(iv) 'ReLU':

- 'ReLU' stands for "Rectified Linear Unit." It is an activation
→ function commonly used in neural networks. The ReLU activation
→ function introduces non-linearity by replacing all negative
→ values in the input with zero and leaving positive values
→ unchanged. Mathematically, it is defined as ' $f(x) = \max(0, x)$ '.
→ ReLU helps the network learn complex patterns and is
→ computationally efficient.

(v) 'linear':

- 'linear' is a common term used in the context of linear regression.
→ However, in PyTorch, the term "linear" is often used to refer to
→ the fully connected layer or linear transformation discussed in (
→ iii) above. It represents a linear mapping from the input to the
→ output, where each input neuron is connected to each output
→ neuron with learned weights.

These functions are essential building blocks for creating and training
→ neural networks in PyTorch, and they play a crucial role in various
→ deep learning tasks, especially in tasks related to image processing
→ and classification.

"""

"""(c) Plot the loss function."""

```
# Plot the training loss
plt.plot(train_losses, label='Training Loss')
plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.title('Training Loss Over Iterations')
plt.legend()
plt.show()
```

"""(d) Edit the code to modify the CNN architecture in the following four
→ steps

(call it myCNN):

"""

```
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
import matplotlib.pyplot as plt

# Define the modified CNN architecture (MyCNN)
```



```

244 class MyCNN(nn.Module):
245     def __init__(self):
246         super(MyCNN, self).__init__()
247         self.conv1 = nn.Conv2d(3, 5, 5)  # (i) 5 activation maps instead
           ↪ of 6
248         self.pool = nn.AvgPool2d(2, 2)  # (ii) Average pooling instead of
           ↪ max pooling
249         self.conv2 = nn.Conv2d(5, 10, 5)  # (iii) 10 activation maps
           ↪ instead of 16
250         self.fc1 = nn.Linear(10 * 5 * 5, 100)  # (iv) 100 dimensions
           ↪ instead of 120
251         self.fc3 = nn.Linear(100, 10)  # Output layer
252
253     def forward(self, x):
254         x = self.pool(nn.functional.relu(self.conv1(x)))
255         x = self.pool(nn.functional.relu(self.conv2(x)))
256         x = x.view(-1, 10 * 5 * 5)
257         x = nn.functional.relu(self.fc1(x))
258         x = self.fc3(x)
259         return x
260
261 # Load CIFAR-10 dataset and create data loaders
262 transform = transforms.Compose(
263     [transforms.ToTensor(),
264      transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
265
266 batch_size = 4
267
268 trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
269                                         download=True, transform=transform
           ↪ )
270 trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
271                                           shuffle=True, num_workers=2)
272
273 # Define the neural network (MyCNN), loss function, and optimizer
274 my_cnn = MyCNN()
275 criterion = nn.CrossEntropyLoss()
276 optimizer = optim.SGD(my_cnn.parameters(), lr=0.001, momentum=0.9)
277
278 # Lists to store training loss values
279 train_losses_my_cnn = []
280
281 # Training loop for MyCNN
282 num_epochs = 2
283 for epoch in range(num_epochs):
284     running_loss = 0.0
285     for i, data in enumerate(trainloader, 0):
286         inputs, labels = data
287         optimizer.zero_grad()
288         outputs = my_cnn(inputs)
289         loss = criterion(outputs, labels)
290         loss.backward()
291         optimizer.step()
292         running_loss += loss.item()

```

```

293         if i % 2000 == 1999: # Print every 2000 mini-batches
294             print(f'[{epoch+1}, {i+1:5d}] loss: {running_loss / \
                ↪ 2000:.3f}')
295             train_losses_my_cnn.append(running_loss / 2000)
296             running_loss = 0.0
297
298 print('Finished Training MyCNN')
299
300 # Plot the training loss for MyCNN
301 plt.plot(train_losses_my_cnn, label='MyCNN Training Loss')
302 plt.xlabel('Iterations (x2000)')
303 plt.ylabel('Loss')
304 plt.title('MyCNN Training Loss Over Iterations')
305 plt.legend()
306 plt.show()
307
308 """To compare the original CNN and the modified MyCNN, we can look at
    ↪ several factors:
309
310 Architecture Differences:
311
312 Original CNN:
313 Conv1: 6 activation maps
314 Max pooling
315 Conv2: 16 activation maps
316 FC1: 120 dimensions
317 FC2: 84 dimensions
318 MyCNN:
319 Conv1: 5 activation maps
320 Average pooling
321 Conv2: 10 activation maps
322 FC1: 100 dimensions
323 FC2 removed
324 Training Loss:
325
326 Compare the training loss curves for both models to see how quickly they
    ↪ converge during training. Lower training loss indicates better
    ↪ convergence.
327 Accuracy on Test Data:
328
329 After training, evaluate both models on the test dataset and compare their
    ↪ accuracy. Higher accuracy indicates better performance.
330 Let's add the evaluation code for the original CNN and then compare the
    ↪ two models:
331 """
332
333 # Evaluate the original CNN on the test dataset
334 correct = 0
335 total = 0
336 with torch.no_grad():
337     for data in testloader:
338         images, labels = data
339         outputs = net(images)
340         _, predicted = torch.max(outputs.data, 1)

```

```

341         total += labels.size(0)
342         correct += (predicted == labels).sum().item()
343
344     print(f'Accuracy of the original CNN on the test images: {100 * correct /
345           ↳ total:.2f}%')
346
347     # Evaluate MyCNN on the test dataset
348     correct_my_cnn = 0
349     total_my_cnn = 0
350     with torch.no_grad():
351         for data in testloader:
352             images, labels = data
353             outputs = my_cnn(images)
354             _, predicted = torch.max(outputs.data, 1)
355             total_my_cnn += labels.size(0)
356             correct_my_cnn += (predicted == labels).sum().item()
357
358     print(f'Accuracy of MyCNN on the test images: {100 * correct_my_cnn /
359           ↳ total_my_cnn:.2f}%')
360
361     """Accuracy of the original CNN on the test images: 53.47%
362     Accuracy of MyCNN on the test images: 52.56%
363     We can see that accuracy is less than in MyCNN compare to Original CNN on
364     ↳ test images.
365     """

```

```

1  [1, 2000] loss: 1.264
2  [1, 4000] loss: 1.235
3  [1, 6000] loss: 1.245
4  [1, 8000] loss: 1.226
5  [1, 10000] loss: 1.248
6  [1, 12000] loss: 1.244
7  [2, 2000] loss: 1.250
8  [2, 4000] loss: 1.238
9  [2, 6000] loss: 1.241
10 [2, 8000] loss: 1.239
11 [2, 10000] loss: 1.234
12 [2, 12000] loss: 1.256
13 Finished Training
14
15 GroundTruth:  cat    ship  ship  plane
16
17 <All keys matched successfully>
18
19 Predicted:  cat    ship  ship  ship
20
21 Accuracy of the network on the 10000 test images: 53 %
22
23 Accuracy for class: plane is 48.0 %
24 Accuracy for class: car   is 78.7 %
25 Accuracy for class: bird  is 45.5 %
26 Accuracy for class: cat   is 36.7 %
27 Accuracy for class: deer  is 32.4 %
28 Accuracy for class: dog   is 54.2 %


```

```

29 Accuracy for class: frog is 69.9 %
30 Accuracy for class: horse is 52.7 %
31 Accuracy for class: ship is 70.6 %
32 Accuracy for class: truck is 46.0 %
33
34
35 Files already downloaded and verified
36 [1, 2000] loss: 2.159
37 [1, 4000] loss: 1.846
38 [1, 6000] loss: 1.690
39 [1, 8000] loss: 1.617
40 [1, 10000] loss: 1.569
41 [1, 12000] loss: 1.538
42 [2, 2000] loss: 1.464
43 [2, 4000] loss: 1.459
44 [2, 6000] loss: 1.451
45 [2, 8000] loss: 1.435
46 [2, 10000] loss: 1.399
47 [2, 12000] loss: 1.396
48 Finished Training MyCNN
49
50 Accuracy of the original CNN on the test images: 53.47%
51 Accuracy of MyCNN on the test images: 52.56%

```





1.3 Problem-2

Understand how to use pretrained CNN for extracting features and fine-tuning using the following video tutorials and associated codes: Code link:

1. (A) <https://www.youtube.com/watch?v=15zlr2vJqKc>
2. (B) <https://www.youtube.com/watch?v=8etkVC93yU4>
3. (c) <https://github.com/madsendennis/notebooks/tree/master/pytorch>

Solution 2:

```

1  # -*- coding: utf-8 -*-
2  """PyTorch_Transfer_learning.ipynb
3
4  Automatically generated by Colaboratory.
5
6  Original file is located at
7      https://colab.research.google.com/drive/1
          ↪ d7YhrpOHC19s2tJonqehckJiqVgzR9-p
8
9  CNN Model For Image Recognition
10
11  Modified Code for CNN Imagae Classification
12  """
13
14  # Mount Google Drive
15  from google.colab import drive
16  drive.mount('/content/drive')
17
18  # Install the gdown library
19  !pip install gdown
20
21  # Define the file ID and output directory
22  file_id = '1fPqPl3X63XqoSWJBVQoCptPszFibjYb3'
23  output_dir = '/content/dataset'
24
25  # Download the file
26  !gdown --id $file_id -O /content/dataset.zip
27
28  # Unzip the dataset
29  !unzip /content/dataset.zip -d $output_dir
30
31  # Update the dataset path
32  dataset = '/content/dataset'
33
34  import torch
35  import torchvision
36  from torchvision import datasets, models, transforms
37  import torch.nn as nn
38  import torch.optim as optim
39  from torch.utils.data import DataLoader
40  import time
41  import os
42  import numpy as np
43  import matplotlib.pyplot as plt
44  from PIL import Image
45  from torchsummary import summary
46
47  # Applying Transforms to the Data

```

```

48 image_transforms = {
49     'train': transforms.Compose([
50         transforms.RandomResizedCrop(size=256, scale=(0.8, 1.0)),
51         transforms.RandomRotation(degrees=15),
52         transforms.RandomHorizontalFlip(),
53         transforms.CenterCrop(size=224),
54         transforms.ToTensor(),
55         transforms.Normalize([0.485, 0.456, 0.406],
56                               [0.229, 0.224, 0.225])
57     ]),
58     'valid': transforms.Compose([
59         transforms.Resize(size=256),
60         transforms.CenterCrop(size=224),
61         transforms.ToTensor(),
62         transforms.Normalize([0.485, 0.456, 0.406],
63                               [0.229, 0.224, 0.225])
64     ]),
65     'test': transforms.Compose([
66         transforms.Resize(size=256),
67         transforms.CenterCrop(size=224),
68         transforms.ToTensor(),
69         transforms.Normalize([0.485, 0.456, 0.406],
70                               [0.229, 0.224, 0.225])
71     ])
72 }
73
74 # Set train and valid directory paths
75 dataset = '/content/dataset/data/'
76 train_directory = os.path.join(dataset, 'train')
77 valid_directory = os.path.join(dataset, 'valid')
78
79 # Batch size
80 bs = 32
81
82 # Number of classes
83 num_classes = len(os.listdir(valid_directory))
84
85 # Load Data from folders
86 data = {
87     'train': datasets.ImageFolder(root=train_directory, transform=
88         ↪ image_transforms['train']),
89     'valid': datasets.ImageFolder(root=valid_directory, transform=
90         ↪ image_transforms['valid'])
91 }
92
93 # Get a mapping of the indices to the class names
94 idx_to_class = {v: k for k, v in data['train'].class_to_idx.items()}
95
96 # Size of Data
97 train_data_size = len(data['train'])
98 valid_data_size = len(data['valid'])
99
100 # Create data loaders
101 train_data_loader = DataLoader(data['train'], batch_size=bs, shuffle=True)

```

```
100 valid_data_loader = DataLoader(data['valid'], batch_size=bs, shuffle=True)
101
102 # Load pre-trained AlexNet model
103 alexnet = models.alexnet(pretrained=True)
104
105 # Freeze model parameters
106 for param in alexnet.parameters():
107     param.requires_grad = False
108
109 # Modify the final layer of AlexNet Model for Transfer Learning
110 alexnet.classifier[6] = nn.Linear(4096, num_classes)
111 alexnet.classifier.add_module("7", nn.LogSoftmax(dim=1))
112
113 # Define Optimizer and Loss Function
114 loss_func = nn.NLLLoss()
115 optimizer = optim.Adam(alexnet.parameters())
116
117 # Define the device (GPU or CPU)
118 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
119
120 # Step 7: Train and Validate the Model
121
122 # Define function to train and validate
123 def train_and_validate(model, loss_criterion, optimizer, epochs=5):
124     start = time.time()
125     history = []
126     best_acc = 0.0
127
128     for epoch in range(epochs):
129         epoch_start = time.time()
130         print("Epoch: {}/{}".format(epoch + 1, epochs))
131
132         # Set to training mode
133         model.train()
134
135         # Loss and Accuracy within the epoch
136         train_loss = 0.0
137         train_acc = 0.0
138
139         valid_loss = 0.0
140         valid_acc = 0.0
141
142         for i, (inputs, labels) in enumerate(train_data_loader):
143             inputs = inputs.to(device)
144             labels = labels.to(device)
145
146             # Clean existing gradients
147             optimizer.zero_grad()
148
149             # Forward pass - compute outputs on input data using the model
150             outputs = model(inputs)
151
152             # Compute loss
153             loss = loss_criterion(outputs, labels)
```



```
154
155     # Backpropagate the gradients
156     loss.backward()
157
158     # Update the parameters
159     optimizer.step()
160
161     # Compute the total loss for the batch and add it to
162         ↪ train_loss
163     train_loss += loss.item() * inputs.size(0)
164
165     # Compute the accuracy
166     ret, predictions = torch.max(outputs.data, 1)
167     correct_counts = predictions.eq(labels.data.view_as(
168         ↪ predictions))
169
170     # Convert correct_counts to float and then compute the mean
171     acc = torch.mean(correct_counts.type(torch.FloatTensor))
172
173     # Compute total accuracy in the whole batch and add to
174         ↪ train_acc
175     train_acc += acc.item() * inputs.size(0)
176
177 # Validation - No gradient tracking needed
178 with torch.no_grad():
179
180     # Set to evaluation mode
181     model.eval()
182
183     # Validation loop
184     for j, (inputs, labels) in enumerate(valid_data_loader):
185         inputs = inputs.to(device)
186         labels = labels.to(device)
187
188         # Forward pass - compute outputs on input data using the
189             ↪ model
190         outputs = model(inputs)
191
192         # Compute loss
193         loss = loss_criterion(outputs, labels)
194
195         # Compute the total loss for the batch and add it to
196             ↪ valid_loss
197         valid_loss += loss.item() * inputs.size(0)
198
199         # Calculate validation accuracy
200         ret, predictions = torch.max(outputs.data, 1)
201         correct_counts = predictions.eq(labels.data.view_as(
202             ↪ predictions))
203
204         # Convert correct_counts to float and then compute the
205             ↪ mean
206         acc = torch.mean(correct_counts.type(torch.FloatTensor))
```

```

201         # Compute total accuracy in the whole batch and add to
           ↪ valid_acc
202         valid_acc += acc.item() * inputs.size(0)
203
204     # Find average training loss and training accuracy
205     avg_train_loss = train_loss / train_data_size
206     avg_train_acc = train_acc / train_data_size
207
208     # Find average training loss and training accuracy
209     avg_valid_loss = valid_loss / valid_data_size
210     avg_valid_acc = valid_acc / valid_data_size
211
212     history.append([avg_train_loss, avg_valid_loss, avg_train_acc,
           ↪ avg_valid_acc])
213
214     epoch_end = time.time()
215
216     print("Epoch: {03d}, Training: Loss: {:.4f}, Accuracy: {:.4f}%,
           ↪ \n\t\tValidation: Loss: {:.4f}, Accuracy: {:.4f}%, Time:
           ↪ {:.4f}s".format(epoch + 1, avg_train_loss, avg_train_acc *
           ↪ 100, avg_valid_loss, avg_valid_acc * 100, epoch_end -
           ↪ epoch_start))
217
218     # Save if the model has the best validation accuracy till now
219     if avg_valid_acc > best_acc:
220         best_acc = avg_valid_acc
221         torch.save(model, dataset + '_model.pt')
222
223     return model, history
224
225 # Specify the number of epochs
226 num_epochs = 5
227
228 # Train and validate the model
229 trained_model, history = train_and_validate(alexnet, loss_func, optimizer,
           ↪ num_epochs)
230
231 # Save training history
232 torch.save(history, dataset + '_history.pt')
233
234 # Print training and validation curves
235 history = np.array(history)
236 plt.plot(history[:, 0], label='Train Loss')
237 plt.plot(history[:, 1], label='Validation Loss')
238 plt.legend()
239 plt.xlabel('Epoch')
240 plt.ylabel('Loss')
241 plt.title('Training and Validation Loss Curves')
242 plt.show()
243
244 plt.plot(history[:, 2], label='Train Accuracy')
245 plt.plot(history[:, 3], label='Validation Accuracy')
246 plt.legend()
247 plt.xlabel('Epoch')

```

```

248 plt.ylabel('Accuracy')
249 plt.title('Training_and_Validation_Accuracy_Curves')
250 plt.show()
251
252 # Define a function to predict the class of a single test image
253 def predict(model, test_image_name, topk=3):
254     transform = image_transforms['test']
255     test_image = Image.open(test_image_name)
256     plt.imshow(test_image)
257     test_image_tensor = transform(test_image).unsqueeze(0)
258
259     if torch.cuda.is_available():
260         test_image_tensor = test_image_tensor.cuda()
261
262     with torch.no_grad():
263         model.eval()
264         out = model(test_image_tensor)
265         ps = torch.exp(out)
266
267         # Check the number of available classes
268         num_classes = ps.shape[1]
269
270         # Adjust topk if there are fewer classes than requested
271         topk = min(topk, num_classes)
272
273         topk_values, topk_indices = ps.topk(topk, dim=1)
274         predictions = []
275
276         for i in range(topk):
277             class_index = topk_indices[0][i].item()
278             class_name = idx_to_class[class_index]
279             score = topk_values[0][i].item()
280             predictions.append((class_name, score))
281
282     return predictions
283
284 # Example usage of the predict function
285 test_image_path = '/content/dataset/data/test/bird/38457.png' # Replace
    ↳ with the path to your test image
286 predictions = predict(trained_model, test_image_path)
287
288 print("Predictions_for", test_image_path)
289 for i, (class_name, score) in enumerate(predictions, start=1):
290     print(f"Prediction_{i}:_Class:_{class_name},_Score:_{score:.4f}")


```



```

1 Epoch: 1/5
2 Epoch : 001, Training: Loss: 0.3292, Accuracy: 85.1190%,
3     Validation : Loss : 0.1369, Accuracy: 94.0909%, Time: 137.5592s
4 Epoch: 2/5
5 Epoch : 002, Training: Loss: 0.2576, Accuracy: 88.5714%,
6     Validation : Loss : 0.1421, Accuracy: 93.1818%, Time: 134.5587s
7 Epoch: 3/5
8 Epoch : 003, Training: Loss: 0.2446, Accuracy: 89.7917%,
9     Validation : Loss : 0.1166, Accuracy: 95.0000%, Time: 136.8005s
10 Epoch: 4/5

```

```
11 Epoch : 004, Training: Loss: 0.2304, Accuracy: 90.1190%,
12     Validation : Loss : 0.1938, Accuracy: 91.8182%, Time: 134.1123s
13 Epoch: 5/5
14 Epoch : 005, Training: Loss: 0.2339, Accuracy: 89.9405%,
15     Validation : Loss : 0.1158, Accuracy: 94.5455%, Time: 139.8286s
16
17
18 Predictions for /content/dataset/data/test/bird/38457.png
19 Prediction 1: Class: bird, Score: 0.9914
20 Prediction 2: Class: horse, Score: 0.0086
```





॥ त्वं ज्ञानमयो विद्वानमयोऽसि ॥